

Project - Image Classification

Importing libraries

```
In [36]: import os
from skimage.io import imread
from skimage.transform import resize
from PIL import Image
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import *
import warnings
warnings.filterwarnings('ignore')
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
import pandas as pd
```

Displaying an image

```
In [9]: Image.open(r"C:\Users\Roy\Python files\Submissions\Project - Image-Classification\cats_and_dogs_filtered\cats_and_dogs_filtered\train\cats\cat.0.jpg")
```



Converting all the images into a 1-D array and storing it into data variable.

Here Label contains two categories i.e 0 (Cat) and 1 (Dog)

```
In [14]: data = []
labels = []

for category_idx, category in enumerate(categories):
    for file in os.listdir(os.path.join(input_dir, category)):
        #loading data
        img_path = os.path.join(input_dir, category, file)
        img = imread(img_path)
        #resizing
        img = resize(img, (15,5))
        data.append(img.flatten())
        labels.append(category_idx)
```

```
In [18]: labels = np.asarray(labels)
data = np.asarray(data)
```

Logistic Regression

```
In [20]: X_train, X_test, y_train, y_test = train_test_split(data, labels, train_size=0.8, shuffle=True, stratify=labels)
```

```
In [21]: sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
In [24]: lr = LogisticRegression()
lr.fit(X_train, y_train)
```

```
Out[24]: LogisticRegression
LogisticRegression()
```

```
In [25]: y_pred = lr.predict(X_test)
```

```
In [26]: accuracy_score(y_test, y_pred)
```

```
Out[26]: 0.6325
```

Decision Tree

```
In [28]: dt = DecisionTreeClassifier()
dt.fit(X_train, y_train)
```

```
Out[28]: DecisionTreeClassifier
DecisionTreeClassifier()
```

```
In [29]: y_pred = dt.predict(X_test)
```

```
In [30]: accuracy_score(y_test, y_pred)
```

```
Out[30]: 0.545
```

Random Forest

```
In [32]: rfc = RandomForestClassifier(n_estimators=1000)
rfc.fit(X_train, y_train)
```

```
Out[32]: RandomForestClassifier
RandomForestClassifier(n_estimators=1000)
```

```
In [33]: y_pred = rfc.predict(X_test)
accuracy_score(y_test, y_pred)
```

```
Out[33]: 0.65
```

Hyperparameter Tuning

```
In [35]: rfc = RandomForestClassifier(random_state=42)
param_grid = {
    'n_estimators': [200, 500],
    'max_features': ['auto', 'sqrt', 'log2'],
    'max_depth': [4, 5, 6, 7, 8],
    'criterion': ['gini', 'entropy']
}
CV_rfc = GridSearchCV(estimator=rfc, param_grid=param_grid, cv=5)
CV_rfc.fit(X_train, y_train)
CV_rfc.best_params_
rfc1 = RandomForestClassifier(random_state=42, max_features='auto', n_estimators=200, max_depth=8, criterion='gini')
rfc1.fit(X_train, y_train)
pred = rfc1.predict(X_test)
print('Accuracy score: ', accuracy_score(y_test, pred))
```

Accuracy score: 0.635

```
In [37]: pd.DataFrame({"Actual Value": y_test, "Predicted Value": y_pred})
```

Out[37]:

	Actual Value	Predicted Value
0	0	0
1	1	0
2	0	1
3	0	1
4	1	0
...
395	1	1
396	0	0
397	1	1
398	0	1
399	1	1

400 rows × 2 columns

```
In [ ]:
```